



UNIVERSITY OF AMSTERDAM
Faculty of Science

Categorical Normalizing Flows via Continuous Transformations

Master Thesis Defense, 25. August 2020

Student: Phillip Lippe



Supervisor: Dr. Efstratios Gavves

Second reader: Dr. Max Welling

Introduction

Motivation


Categorical Data



Introduction

Preliminaries

Normalizing Flows




- + Universality
- + Exact likelihood estimate
- + Efficient density evaluation and (parallel) sampling

Figure credit: Weng, Lilian. "Flow-based Deep Generative Models", 2018.

Introduction

Related Work


Applying (Variational) Dequantization



Designed for image modeling

- Categories are not “quantized” real values

Discrete Normalizing Flows



- Is limited to permutations
- Not universal with factorized prior
- (Biased) gradient approximations and difficult optimization

References


Tran, D. et al.: “Discrete Flows: Invertible Generative Models of Discrete Data”. NeurIPS, 2019.
Hoogeboom, E. et al.: “Integer Discrete Flows and Lossless Compression”. NeurIPS, 2019.

Introduction

Contributions

Categorical Normalizing Flow



Modeling categorical distribution by a continuous normalizing flow



- + Universality
- + Stable optimization without biased gradients
- + Efficient density evaluation and (parallel) sampling

GraphCNF

Powerful graph generation model based on Categorical Normalizing Flows



- + One-shot generation
- + Permutation-invariant to node order
- + Support of categorical node and edge attributes

Categorical Normalizing Flows

Encoding

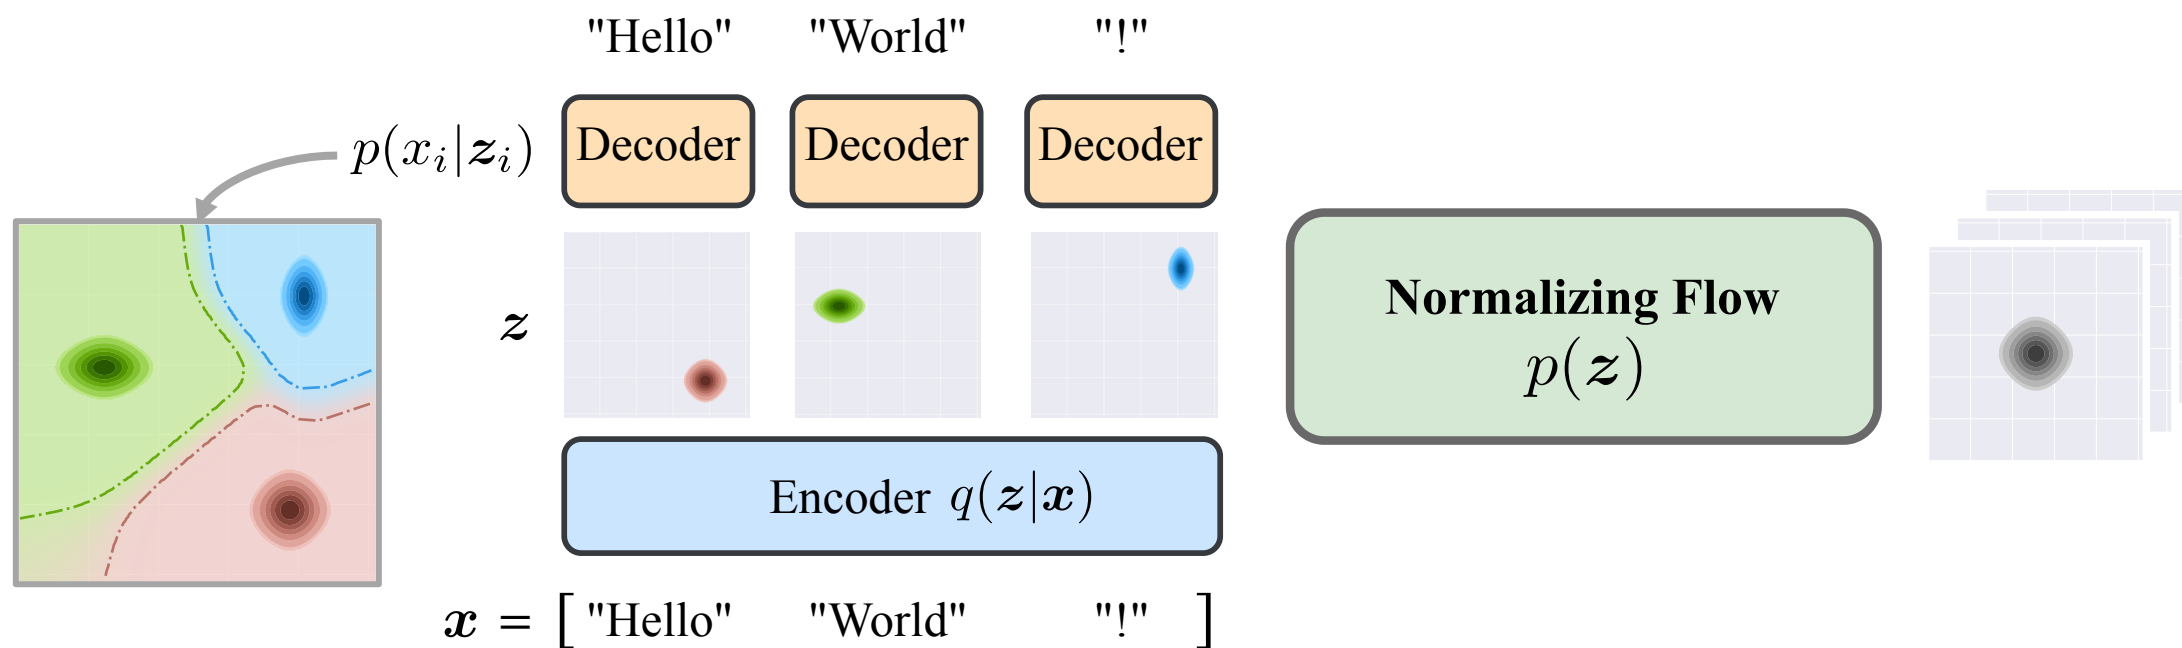
- First step: represent categorical data in continuous space
- Desired properties of an encoding function
 - No loss of information (non-overlapping volumes)
 - Learnable
 - Smooth
 - Support for higher dimensions

⇒ Variational inference with factorized posterior: $p(\mathbf{x}) \geq \mathbb{E}_{\mathbf{z} \sim q(\cdot|\mathbf{x})} \left[\frac{\prod_i p(x_i|\mathbf{z}_i)}{q(\mathbf{z}|\mathbf{x})} p(\mathbf{z}) \right]$

- Ensures that continuous form \mathbf{z} contains the exact same information as discrete \mathbf{x}
- Small variational gap \Rightarrow close-to exact likelihood estimate

Categorical Normalizing Flows

Overview




Objective function:
$$p(\mathbf{x}) \geq \mathbb{E}_{\mathbf{z} \sim q(\cdot|\mathbf{x})} \left[\frac{\prod_i p(x_i|z_i)}{q(\mathbf{z}|\mathbf{x})} p(\mathbf{z}) \right]$$

Categorical Normalizing Flows

Encoding function

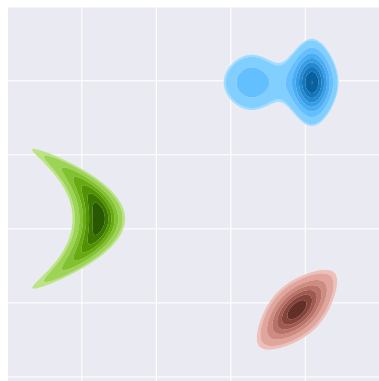
Mixture model



$$q(\mathbf{z}|\mathbf{x}) = \prod_{i=1}^N g(\mathbf{z}_i|\mu(x_i), \sigma(x_i))$$

- Exact posterior can be found
⇒ no “variational” gap

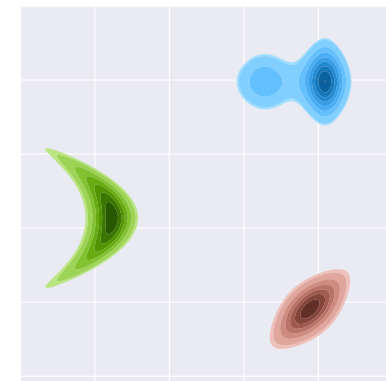
Linear flows



$$q(\mathbf{z}|\mathbf{x}) = \prod_{i=1}^N q(\mathbf{z}_i|x_i)$$

- One flow per category
- Shared across input

Variational encoding




$$q(\mathbf{z}|\mathbf{x})$$

- One flow on the whole discrete input
- Similar to variational dequantization

Categorical Normalizing Flows

Related Work – Flow-based Variational Autoencoders



- Non-factorized posterior
- Complexity split between decoder and normalizing flow
- Instable training
- Performance worse than non-latent baseline

References

Ziegler, Z.M. and Rush, A.M.: “Latent Normalizing Flows for Discrete Sequences”. ICML, 2019.

Categorical Normalizing Flows


Experiments – Set Modeling

- Toy datasets on sets with known dataset likelihood
- **Metric:** test likelihood in bits per categorical variable (lower = better)

Model	Set shuffling	Set summation
Discrete NF	3.87 \pm 0.04	2.51 \pm 0.00
Variational Dequantization	3.01 \pm 0.02	2.29 \pm 0.01
Latent NF	2.78 \pm 0.00	2.26 \pm 0.01
CNF + Mixture model	2.78 \pm 0.00	2.24 \pm 0.00
CNF + Linear flows	2.78 \pm 0.00	2.25 \pm 0.00
CNF + Variational Encoding	2.79 \pm 0.01	2.25 \pm 0.01
Optimal	2.77	2.24

Categorical Normalizing Flows

Experiments – Language Modeling




Categorical Normalizing Flows

Conclusion

- Categorical Normalizing Flow is a powerful framework to accurately model categorical distributions
- Encoding function: Mixture model efficient with exact posterior
- Factorized posterior pushes all complexity into the prior

Graph Generation with CNF

Introduction



(1) Node attributes 

(2) Edge attributes 


(3) Adjacency matrix 

Challenge: nodes are unordered, i.e. a set
⇒ Maintain permutation-invariance of nodes


Graph Generation with CNF

GraphCNF


CNF - Node type representation
(discrete \rightarrow continuous)




CNF - Edge attribute representation
(discrete \rightarrow continuous)



Adding virtual edge
representation
(CNF)



Prior distribution



- + Permutation-invariant
- + Efficient three-step approach


Graph Generation with CNF

Experiments – Molecule Generation

- **Task:** given a set of molecules, learn to model the space of valid molecules
- **Metrics:** calculated on 10k generated graphs,
 - (1) *Validity*: percentage of graphs being valid molecules
 - (2) *Uniqueness*: percentage of unique molecules
 - (3) *Novelty*: percentage of molecules that are not equal to any training molecule
 - (4) *Reconstruction*: reconstruction accuracy of test molecules from latent space

Graph Generation with CNF

Experiments – Molecule Generation







Results on the Zinc250k dataset (224k examples)

Method	Validity	Uniqueness	Novelty	Reconstruction	Parallel	General
JT-VAE	100%	100%	100%	71%	✗	✗
GraphAF	68%	99.10%	100%	100%	✗	✓
R-VAE	34.9%	100%	—	54.7%	✓	✓
GraphNVP	42.60%	94.80%	100%	100%	✓	✓
GraphCNF	83.41%	99.99%	100%	100%	✓	✓
	(± 2.88)	(± 0.01)	(± 0.00)	(± 0.00)		
+ Sub-graphs	96.35%	99.98%	99.98%	100%	✓	✓
	(± 2.21)	(± 0.01)	(± 0.02)	(± 0.00)		

Graph Generation with CNF

Experiments – Molecule Generation



Conclusion

- Mixture model encoding is the “dequantization” for categorical data
 - Simple, efficient and exact in likelihood (continuous)
- CNFs enable strong, latent-based generative models on domains like graphs
 - GraphCNF significantly outperforms previous flow-based approach on molecule generation
- Possible future direction:
 - Combining continuous and discrete normalizing flows
 - GraphCNF on large graphs ($|V| > 100$)

Thank You